

# Hands-on Guide: Using Valkey with Amazon ElastiCache Serverless



## **Table of Contents**

1. Introduction
  2. What is Valkey?
  3. Why Migrate from Redis to Valkey?
  4. Setting Up Valkey with Amazon ElastiCache Serverless
  5. Connecting to Valkey
  6. Programmatic Access Using Valkey Clients
  7. Benchmarking Valkey vs Redis OSS
  8. Seamless Upgrade from Redis OSS to Valkey
  9. Summary
  10. Best Practices
-

## Introduction

As demand for high-performance caching and messaging systems grows, developers and businesses are exploring open-source alternatives to Redis. One such emerging solution is Valkey, a Redis-compatible, community-driven, open-source key-value store.

This guide is a comprehensive reference for:

- Deploying Valkey using Amazon ElastiCache Serverless
  - Performing performance benchmarking
  - Seamlessly upgrading from Redis OSS to Valkey
- 

## What is Valkey?



Valkey is a high-performance, in-memory key-value data store designed for fast, scalable, and efficient data operations. It maintains Redis protocol compatibility, allowing it to act as a drop-in replacement.

### Core Use Cases:

- In-memory data caching
- Session management
- Message queues
- Real-time analytics
- Primary NoSQL database for low-latency applications

## Community and Licensing:

- Backed by the Linux Foundation
  - BSD-licensed
  - 21,000+ GitHub stars
  - 800+ contributors
  - Actively maintained and improved
- 

## Why Migrate from Redis to Valkey?

Valkey is quickly becoming the preferred Redis alternative due to the following reasons:

### 1. Licensing Flexibility

Redis transitioned from an open-source license to AGPL, introducing restrictions for commercial use. Valkey remains BSD-licensed, making it safe and flexible for enterprise deployment.

### 2. Open Governance

Valkey is governed by an open community under the Linux Foundation, ensuring transparent development and long-term reliability.

### 3. Performance Enhancements

Valkey delivers improved throughput and reduced latency over Redis OSS. Optimizations include lower CPU usage and memory overhead.

### 4. Compatibility

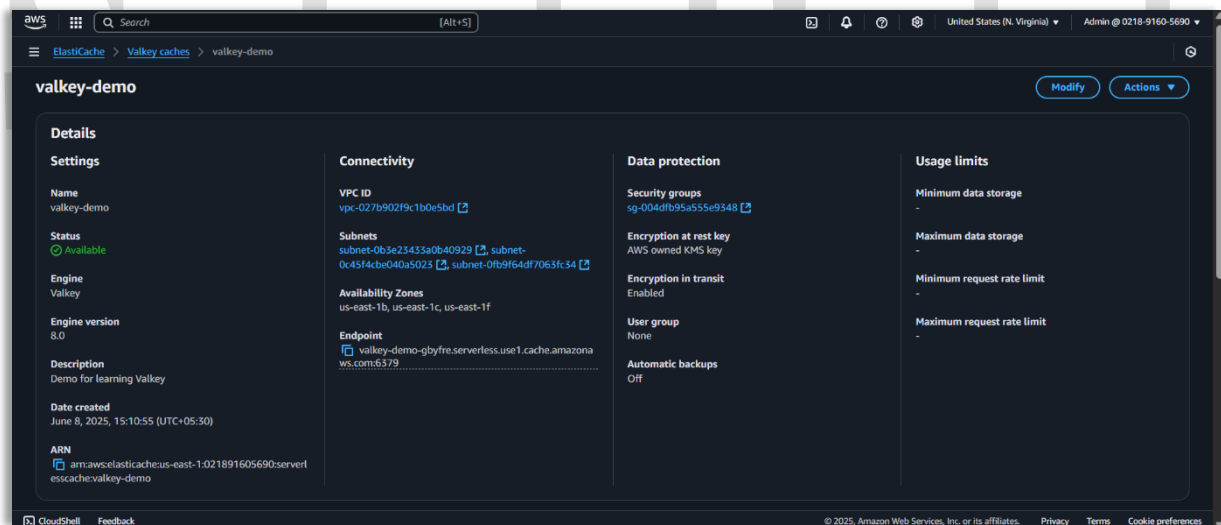
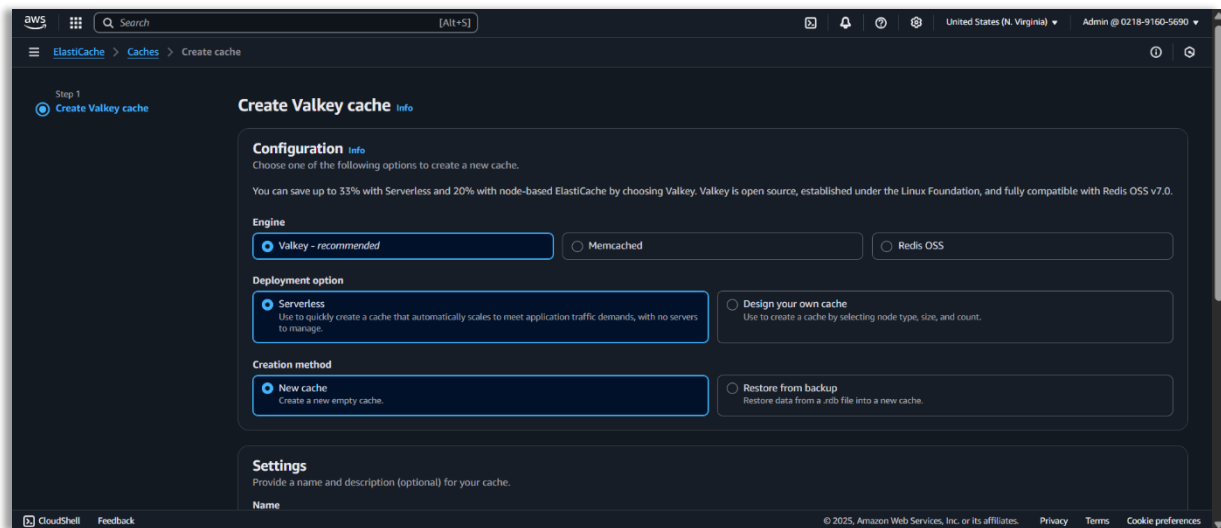
Valkey supports the same protocol and commands as Redis, making it a drop-in replacement with minimal application-level changes.

---

## Setting Up Valkey with Amazon ElastiCache Serverless

Amazon provides Valkey as a serverless managed option in ElastiCache, allowing developers to focus on application logic instead of infrastructure.

### Setup Steps:



1. Log in to the AWS Console
2. Navigate to Amazon ElastiCache
3. Click "Get Started"
4. Choose "Valkey" as the engine
5. Select "Serverless" deployment

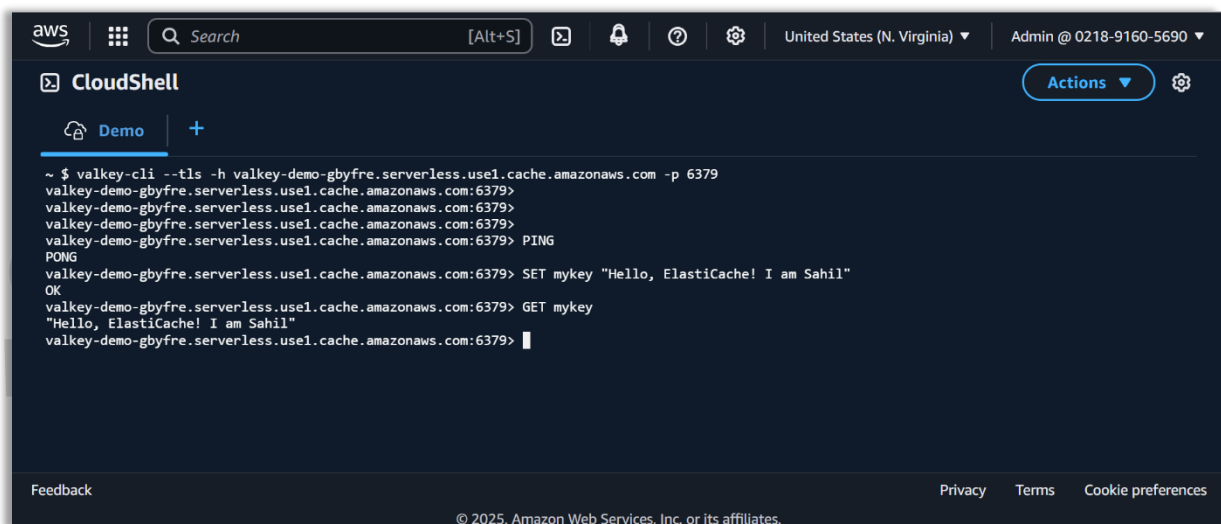
6. Fill in configuration:
  - Name: valkey-demo
  - Engine Version: 8 (recommended)
7. Adjust advanced settings (VPC, encryption, etc.)
8. Click "Create"

In a few minutes, your cache becomes active and provides an endpoint for connectivity.

---

## Connecting to Valkey

### Using AWS CloudShell:



The screenshot shows the AWS CloudShell interface. The terminal window has a title bar with 'aws', a search bar, and navigation icons. Below the title bar, there's a 'CloudShell' header with a 'Demo' tab and an 'Actions' button. The terminal content shows the following commands and output:

```
~ $ valkey-cli --tls -h valkey-demo-gbyfre.serverless.us1.cache.amazonaws.com -p 6379
valkey-demo-gbyfre.serverless.us1.cache.amazonaws.com:6379>
valkey-demo-gbyfre.serverless.us1.cache.amazonaws.com:6379>
valkey-demo-gbyfre.serverless.us1.cache.amazonaws.com:6379>
valkey-demo-gbyfre.serverless.us1.cache.amazonaws.com:6379> PING
PONG
valkey-demo-gbyfre.serverless.us1.cache.amazonaws.com:6379> SET mykey "Hello, ElastiCache! I am Sahil"
OK
valkey-demo-gbyfre.serverless.us1.cache.amazonaws.com:6379> GET mykey
"Hello, ElastiCache! I am Sahil"
valkey-demo-gbyfre.serverless.us1.cache.amazonaws.com:6379> █
```

At the bottom of the terminal window, there are links for 'Feedback', 'Privacy', 'Terms', and 'Cookie preferences', along with a copyright notice: '© 2025, Amazon Web Services, Inc. or its affiliates.'

1. From the cache instance page, go to "Connectivity & Security"
2. Click "Connect to cache" to launch AWS CloudShell
3. Run the following command to connect:
4. `./valkey-cli -h <your-endpoint> -p 6379`
5. Test with basic commands:
6. set greeting "Hello, Valkey"
7. get greeting

Expected output:

"Hello, Valkey"

---

## Programmatic Access Using Valkey Clients

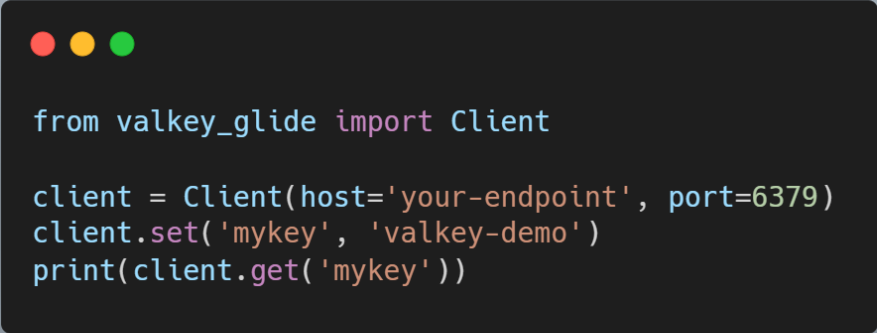
You can integrate Valkey into your applications using Valkey Glide or other Redis-compatible libraries.

### Python Example:

Install the client:

pip install valkey-glide

Sample script:



```
from valkey_glide import Client

client = Client(host='your-endpoint', port=6379)
client.set('mykey', 'valkey-demo')
print(client.get('mykey'))
```

Valkey also works with existing Redis libraries like redis-py, ioredis, or Jedis.

---

## Benchmarking Valkey vs Redis OSS

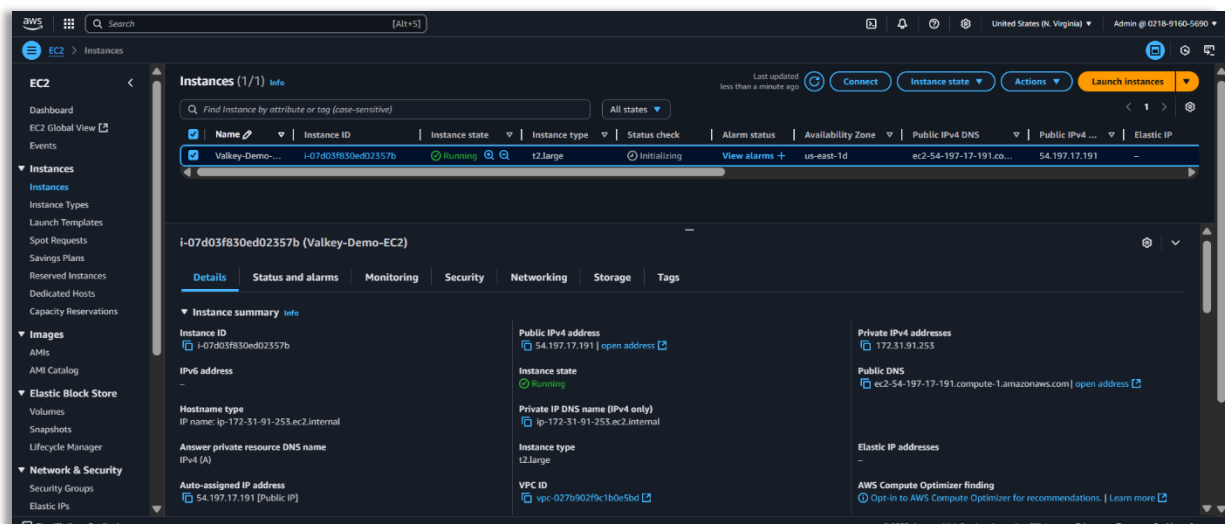
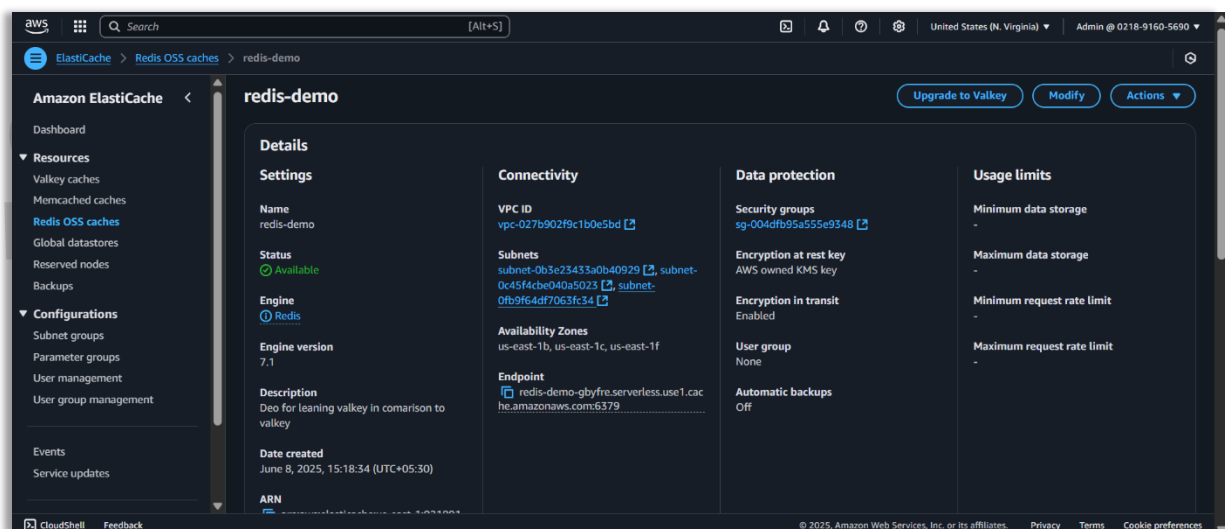
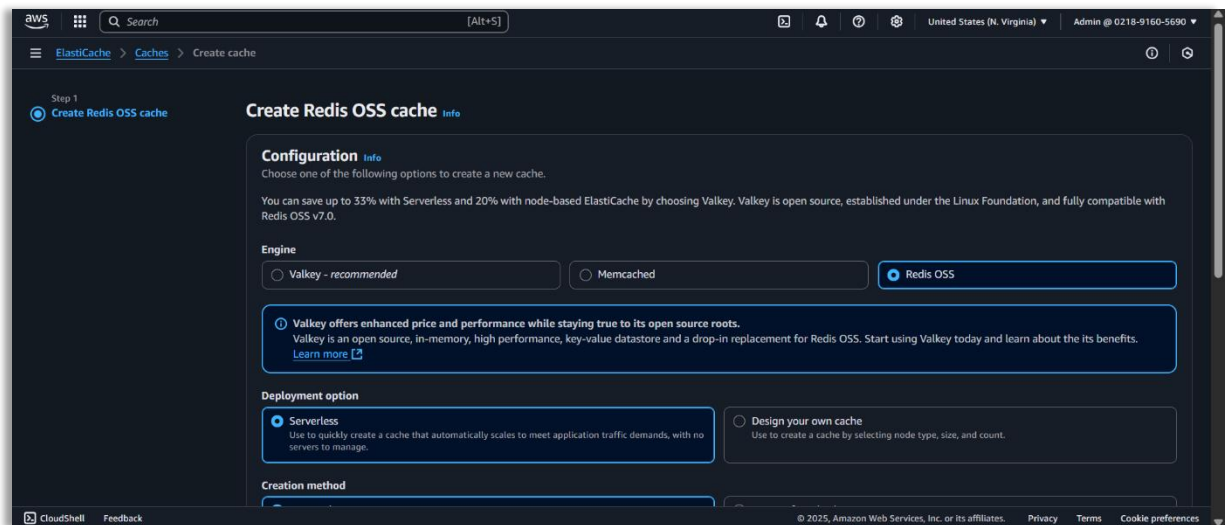
### Objective:

Evaluate performance across key metrics:

- Operations per second (throughput)
- Latency (P50, P90, P99)

## Setup:

- Two ElastiCache instances: one Valkey, one Redis OSS, EC2 instance to run benchmarking script, Python script performing set, get, and delete operations



```

import redis
import time
import random
import string
import statistics

# Redis connection details
REDIS_HOST = "redisval-bu8xv7.serverless.usel.cache.amazonaws.com"
REDIS_PORT = 6379 # TLS port
REDIS_PASSWORD = None

NUM_OPERATIONS = 10000
KEY_PREFIX = "bench_key_"

def random_string(length=10):
    return ''.join(random.choices(string.ascii_letters + string.digits,
    k=length))

def benchmark_operation(redis_client, op_name, action_fn):
    latencies = []
    start = time.time()

    for _ in range(NUM_OPERATIONS):
        t0 = time.perf_counter()
        action_fn()
        t1 = time.perf_counter()
        latencies.append((t1 - t0) * 1000) # latency in ms

    total_time = time.time() - start
    throughput = NUM_OPERATIONS / total_time
    avg_latency = statistics.mean(latencies)
    p50 = statistics.median(latencies)
    p90 = statistics.quantiles(latencies, n=100)[89]
    p99 = statistics.quantiles(latencies, n=100)[98]

    return {
        "op": op_name,
        "total_time": total_time,
        "throughput": throughput,
        "avg_latency_ms": avg_latency,
        "p50_latency_ms": p50,
        "p90_latency_ms": p90,
        "p99_latency_ms": p99,
    }

def benchmark_redis(redis_client):
    results = []
    counter = 0

    def set_fn():
        nonlocal counter
        key = f"{KEY_PREFIX}{counter}"
        value = random_string(50)
        redis_client.set(key, value)
        counter += 1

    results.append(benchmark_operation(redis_client, "SET", set_fn))

    counter = 0
    def get_fn():
        nonlocal counter
        key = f"{KEY_PREFIX}{counter}"
        redis_client.get(key)
        counter += 1

    results.append(benchmark_operation(redis_client, "GET", get_fn))

    counter = 0
    def del_fn():
        nonlocal counter
        key = f"{KEY_PREFIX}{counter}"
        redis_client.delete(key)
        counter += 1

    results.append(benchmark_operation(redis_client, "DEL", del_fn))

    return results

def main():
    try:
        client = redis.Redis(
            host=REDIS_HOST,
            port=REDIS_PORT,
            password=REDIS_PASSWORD,
            ssl=True, # your TLS enabled here
            decode_responses=True
        )

        client.ping()
        print(f"✅ Connected to Redis at {REDIS_HOST}:{REDIS_PORT} over TLS")
        print(f"🚀 Running benchmark with latency stats...")

        results = benchmark_redis(client)

        print("\n📊 Benchmark Results:")
        for r in results:
            print(f"\n💎 {r['op']} Operation")
            print(f"    Total Time      : {r['total_time']:.4f} sec")
            print(f"    Throughput      : {r['throughput']:.2f} ops/sec")
            print(f"    Average Latency: {r['avg_latency_ms']:.3f} ms")
            print(f"    p50 Latency     : {r['p50_latency_ms']:.3f} ms")
            print(f"    p90 Latency     : {r['p90_latency_ms']:.3f} ms")
            print(f"    p99 Latency     : {r['p99_latency_ms']:.3f} ms\n")

    except Exception as e:
        print(f"❌ Error: {e}")

if __name__ == "__main__":
    main()

```



1. Edit benchmarking.py to include your cache endpoint.
2. Run: `python3 benchmarking.py`

### Observations:

```
ubuntu@ip-172-31-91-253: ~  
(venv) ubuntu@ip-172-31-91-253:~$ python3 benchmarking.py  
☑ Connected to Redis at valkey-demo-gbyfre.serverless.use1.cache.amazonaws.com:6379 over TLS  
🔗 Running benchmark with latency stats...  
  
📊 Benchmark Results:  
  
◦ SET Operation  
Total Time      : 16.8429 sec  
Throughput      : 593.72 ops/sec  
Average Latency: 1.683 ms  
p50 Latency     : 1.627 ms  
p90 Latency     : 2.142 ms  
p99 Latency     : 2.634 ms  
  
◦ GET Operation  
Total Time      : 16.3526 sec  
Throughput      : 611.52 ops/sec  
Average Latency: 1.634 ms  
p50 Latency     : 1.584 ms  
p90 Latency     : 2.059 ms  
p99 Latency     : 2.545 ms  
  
◦ DEL Operation  
Total Time      : 16.1707 sec  
Throughput      : 618.40 ops/sec  
Average Latency: 1.616 ms  
p50 Latency     : 1.568 ms  
p90 Latency     : 2.032 ms  
p99 Latency     : 2.539 ms
```

```
ubuntu@ip-172-31-91-253: ~
(venv) ubuntu@ip-172-31-91-253:~$ ls
benchmarking.py  venv
(venv) ubuntu@ip-172-31-91-253:~$python3 benchmarking.py
☑ Connected to Redis at redis-demo-gbyfre.serverless.use1.cache.amazonaws.com:6379 over TLS
🔗 Running benchmark with latency stats...

📊 Benchmark Results:

◦ SET Operation
Total Time      : 22.1257 sec
Throughput      : 451.96 ops/sec
Average Latency: 2.211 ms
p50 Latency     : 2.165 ms
p90 Latency     : 2.609 ms
p99 Latency     : 3.092 ms

◦ GET Operation
Total Time      : 21.8218 sec
Throughput      : 458.26 ops/sec
Average Latency: 2.181 ms
p50 Latency     : 2.133 ms
p90 Latency     : 2.569 ms
p99 Latency     : 3.060 ms

◦ DEL Operation
Total Time      : 21.8263 sec
Throughput      : 458.16 ops/sec
Average Latency: 2.181 ms
p50 Latency     : 2.140 ms
p90 Latency     : 2.573 ms
p99 Latency     : 3.050 ms
```

- Valkey consistently showed lower latency at all percentiles
- Throughput was higher in all tested workloads
- Especially strong in high-concurrency and read-heavy scenarios

---

## Seamless Upgrade from Redis OSS to Valkey

Amazon ElastiCache Serverless allows one-click upgrades from Redis OSS to Valkey.

### Upgrade Steps:

1. Open your Redis OSS cache in ElastiCache
2. Click "Upgrade to Valkey"
3. Select Engine Version 8
4. Confirm and proceed

Upgrade completes in minutes with no data loss or service downtime.

---

## Summary

Feature	Redis OSS	Valkey
License	AGPL	BSD
Governance	Redis Ltd	Linux Foundation
Compatibility	Native Redis	Redis-Compatible
Performance	Good	Better
Cloud Support	Yes	Yes
Upgrade Path	Manual	One-click in AWS

---

## Best Practices

- Benchmark workloads before and after migration
  - Enable encryption in transit and at rest for production use
  - Choose Valkey Glide for advanced client features
  - Test upgrades in a non-production environment
  - Use CloudWatch for metrics and alerting
  - Leverage serverless deployment for cost optimization
-